

## Model

- To extract the embeddings, I chose a pretrained model that was trained on law corpus. This choice was also guided by the fact that such a model may be better equipped to yield contextualised embedding vectors than models trained on unrelated corpus.
- Decided to add an extra flexibility of obtaining either normalised or non-normalised embeddings

## Indexing

- Decided to handle document indexing by bulk streaming

## Retrieval

- Decided to add a fix ( timeout, max\_retries, retry\_on\_timeout ) to handle request time-out errors due to heavy payload in request body caused by long queries / texts in the search response objects

## ElasticSearch setup & Debugging

- Decided to setup elastic search with docker locally
- Although setup was successful, the server unexpectedly failed to communicate with any request after countless attempts and tricks
- Unfortunately I could not debug the issue to generate results and progress to the other tasks that relied on these results